



CBS Extensions

- Non-holonomic agents
- Mix of non-anonymous agents (= designated goal locations) and anonymous agents (= assignable goal locations)
- Agents of different sizes
- Deadlines
- **Uncertainty about the speeds of agents**



see earlier credits

- [1] H. Ma and S. Koenig., "Optimal Target Assignment and Path Finding for Teams of Agents", AAMAS 2016.
 [2] L. Cohen et al., "Optimal and Bounded-Suboptimal Multi-Agent Motion Planning", SoCS 2019.
 [3] J. Li et al., "Multi-Agent Path Finding for Large Agents", AAAI 2019.
 [4] H. Ma et al., "Multi-Agent Path Finding with Deadlines", IJCAI 2018.

1

Robustness to Delays during Planning and Plan Execution




2



Robust Planning with CBS Extensions

- A MAPF plan is k -robust iff no collision occurs even if each agent can be delayed during execution by at most k timesteps



[1] D. Atzmon et al., "k-Robust Multi-Agent Path Finding", SoCS, 2017.

3



Robust Planning with CBS Extensions

- A MAPF plan is p -collision free iff the probability that no collision occurs during its execution is at most p

[1] D. Atzmon et al., "Probabilistic Robust Multi-Agent Path Finding", ICAPS, 2020.

4



Robust Planning with CBS Extensions

- A MAPF plan is p -collision free iff the probability that no collision occurs during its execution is at most p
 - **Detection of potential collisions:**
two agents moving at the nominal speed can now collide even if they are in the same location at different time steps
 - **Resolution of potential collisions:**
add a third possibility, namely that the collision will occur
 - **Verification that a MAPF plan is collision-free:**
Calculate the probability that no collision occurs and ensure that it is at least p

[1] D. Atzmon et al., "Probabilistic Robust Multi-Agent Path Finding", ICAPS, 2020.

5



Robust Plan Execution

- Planning always uses models that are not completely accurate
 - Robots have unmodeled kinematic constraints
 - ...
- Plan execution will therefore deviate from the plan
- Strategy 1: If one robot is delayed, delay all robots
 - Problem: Bad throughput
- **Strategy 2:** Replan whenever plan execution deviates from the plan
 - Problem: Planning is slow since finding good plans is NP-hard


[1] W. Hoenig et al., "Multi-Agent Path Finding with Kinematic Constraints", ICAPS, 2016.

6

Robust Plan Execution: ML

not our work

- PRIMAL: mix of deep reinforcement learning and imitation learning



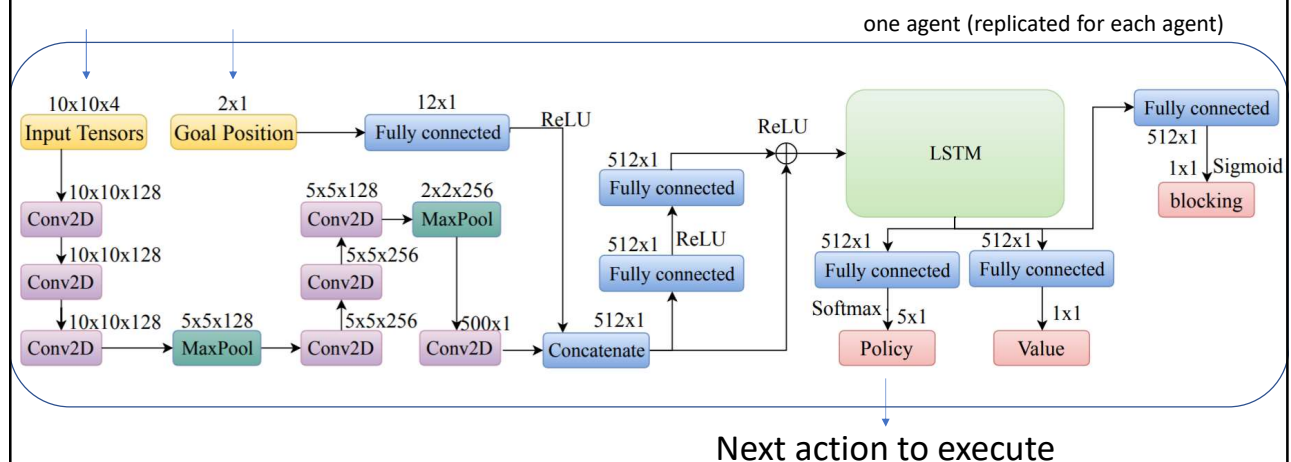
<https://medium.com/free-code-camp/explained-simply-how-deepmind-taught-ai-to-play-video-games-9eb5f38c89ee> <https://www.pcworld.com/article/2889432/google-ai-program-masters-classic-atari-video-games.html>

[1] V. Mnih et al., "Human Level Control through Deep Reinforcement Learning", Nature 518, 2015.

7

Robust Plan Execution: ML

- PRIMAL: mix of deep reinforcement learning and imitation learning




[1] G. Sartoretti et al., "PRIMAL: Pathfinding via Reinforcement and Imitation Multi-Agent Learning," IEEE RAL 4(3), 2019.

8

Robust Plan Execution: ML

- PRIMAL: mix of deep reinforcement learning and imitation learning

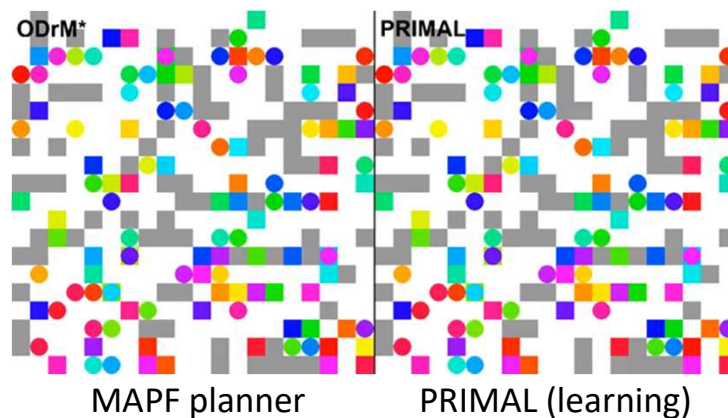


[1] G. Sartoretti et al., "PRIMAL: Pathfinding via Reinforcement and Imitation Multi-Agent Learning," IEEE RAL 4(3), 2019.

9

Robust Plan Execution: ML

- Training: 20 days in a supercomputing center
- Ideal plan execution: 64 agents – 20x20 map – 20% obstacle density




[1] G. Sartoretti et al., "PRIMAL: Pathfinding via Reinforcement and Imitation Multi-Agent Learning," IEEE RAL 4(3), 2019.

11

ML for Planning

- Enhancing planning with machine learning
 - Planners use lots of hard-coded decision strategies
 - Machine learning can often learn to make better decisions
 - The resulting planners can be more efficient and/or effective



[1] T. Huang et al., "Learning to Resolve Conflicts for Multi-Agent Path Finding with Conflict-Based Search," AAAI, 2021.

[2] T. Huang et al., "Learning Node-Selection Strategies in Bounded-Suboptimal Conflict-Based Search for Multi-Agent Path Finding"

12

Robust Plan Execution

- Planning always uses models that are not completely accurate
 - Robots have unmodeled kinematic constraints
 - ...
- Plan execution will therefore deviate from the plan
- Strategy 1: If one robot is delayed, delay all robots
 - Problem: Bad throughput
- Strategy 2: Replan whenever plan execution deviates from the plan
 - Problem: Planning is slow since finding good plans is NP-hard

[1] W. Hoenig et al., "Multi-Agent Path Finding with Kinematic Constraints", ICAPS, 2016.

13



Robust Plan Execution: Reasoning

- Find a MAPF plan (**slow**).
- In a feedback loop, repeatedly determine the speed with which each robot should move along its path given its current position (**fast**).
- Only if the problem becomes unsolvable, find a new MAPF plan.

[1] W. Hoenig et al., "Multi-Agent Path Finding with Kinematic Constraints", ICAPS, 2016.

14




Robust Plan Execution: Reasoning

- MAPF-POST makes use of a simple temporal network to post-process a given MAPF plan in polynomial time to allow for robust plan execution on robots
 - Takes into account edge lengths
 - Takes into account speed limits on edges
 - Takes into account maximum velocities of robots
 - Guarantees safety distances

[1] W. Hoenig et al., "Multi-Agent Path Finding with Kinematic Constraints", ICAPS, 2016.


15

Robust Plan Execution: Reasoning



16

Robust Plan Execution: Reasoning



17

Robust Plan Execution: Reasoning



18


Robust Plan Execution: Reasoning



19




Robust Plan Execution: Reasoning



20




Robust Plan Execution: Reasoning



21

Robust Plan Execution: Reasoning



Agent 1 $A \rightarrow B \rightarrow C \rightarrow D \rightarrow E$


Agent 2 $B \rightarrow C \rightarrow F \rightarrow C \rightarrow D$

Precedence Graph

vertex = event that an agent arrives at a location

Agent 1


Agent 2



[1] W. Hoenig et al., "Multi-Agent Path Finding with Kinematic Constraints", ICAPS, 2016.

22

Robust Plan Execution: Reasoning



Agent 1 $A \rightarrow B \rightarrow C \rightarrow D \rightarrow E$


Agent 2 $B \rightarrow C \rightarrow F \rightarrow C \rightarrow D$

Precedence Graph

Type 1 edge = order in which the same agent arrives at locations

Agent 1

Agent 2



[1] W. Hoenig et al., "Multi-Agent Path Finding with Kinematic Constraints", ICAPS, 2016.

23



Robust Plan Execution: Reasoning

- Minimize makespan and flowtime
 - Determine the earliest arrival times in each location that satisfy the constraints.
 - Calculate the speeds of the robots based on the arrival times.

$$\begin{aligned}
 &\text{Minimize } \sum_{j=1}^K t(v^j) \\
 &\text{such that } t(X_S) = 0 \\
 &\text{and, for all } e = (v, v') \in \mathcal{E}', \\
 &t(v') - t(v) \geq LB(e) \\
 &t(v') - t(v) \leq UB(e)
 \end{aligned}$$

[1] W. Hoenig et al., "Multi-Agent Path Finding with Kinematic Constraints", ICAPS, 2016.

[2] R. Dechter et al., "Temporal Constraint Networks", AIJ 49(1-3), 1991.

26



Robust Plan Execution: Reasoning



[1] W. Hoenig et al., "Multi-Agent Path Finding with Kinematic Constraints", ICAPS, 2016.

28

Robust Plan Execution: Reasoning




[1] W. Hoenig et al., "Multi-Agent Path Finding with Kinematic Constraints", ICAPS, 2016.

29

Robust Plan Execution: Reasoning

not our work

- Mixed reality simulation



[1] W. Hoenig et al., "Mixed Reality for Robots", IROS, 2015.

30